

# Language & Technology

## Lecture 4: More Word-Based Models

Alëna Aksënova & Aniello De Santo

Stony Brook University  
`alena.aksenova@stonybrook.edu`  
`aniello.desanto@stonybrook.edu`

## Previously on “Language & Technology”...

- ▶ Counting words!

`string` = "The sun shone, having no alternative, on the  
nothing new."

`tokens` ["the", "sun", "shone", "having", "no",  
"alternative", "on", "the", "nothing", "new"]

- ▶ More technically: for each word type we count its word tokens.

`word type` a word of a given language

`word token` instance of the word in a text

# Previously on “Language & Technology”...

## Steps of Word Counting Model

- 1 collect corpus
- 2 tokenize strings
- 3 count tokens for each type

Counting words is surprisingly useful:

- ▶ Culturomics
- ▶ Stylistic analysis

But there's more!

# Previously on “Language & Technology”...

## Steps of Word Counting Model

- 1 collect corpus
- 2 tokenize strings
- 3 count tokens for each type

Counting words is surprisingly useful:

- ▶ Culturomics
- ▶ Stylistic analysis

But there's more!

## Application 3: Authorship Attribution

- ▶ Computational stylistic analysis isn't just good at predicting success of novels.
- ▶ Word counts pick up on **statistical tendencies in writing** of individual authors
- ▶ This means that potential authors can be inferred from word frequencies!

# Example: Auditing the Bard's Oeuvre

- ▶ Several works of Shakespeare have supposedly been (partially) authored by somebody else:
  - ▶ *Edward III*  
co-authored with Thomas Kyd
  - ▶ *The Comedy of Errors*,  
*Love's Labour Lost*,  
*The Tempest*  
co-authored with Francis Bacon
- ▶ These are **fringe ideas** in literary circles.
- ▶ But there are word counting models that **support these claims**.



# Another Example: Harry Potter's Mum

- ▶ New book: *The Cuckoo's Calling* by Robert Galbraith
- ▶ *Sunday Times* suspected that to be a pseudonym for J. K. Rowling.
- ▶ Evidence from word-based model:
  - ▶ 100 most common words
  - ▶ average word length
  - ▶ word pairs  
(discussed in next lecture)
  - ▶ character 4-grams  
(discussed in next lecture)



## A Friendly Pointer

Article on LanguageLog:

<http://languagelog.ldc.upenn.edu/n11/?p=5315>

- ▶ As with culturomics, it is unclear what to make of the findings.
  - ▶ J.K. Rowling: success!
  - ▶ Shakespeare: success?
- ▶ One can get vastly different results depending on:
  - ▶ absolute size of test corpus
  - ▶ relative size (e.g. ratio of Shakespeare to Bacon)
  - ▶ whether certain words are ignored
  - ▶ whether one pays attention to rare words or frequent words
- ▶ Many of these studies feel backwards:
  - ▶ They start with a specific hypothesis and then present one specific model that supports this hypothesis.
  - ▶ But why is this model better than all of the alternatives?



# Moral of the Story

- ▶ Authorship attribution models still ill-understood
- ▶ In the future, these models might answer long-standing questions or pose new ones.
- ▶ Maybe you can try it in an English class.

## Application 4: Word Meanings

- ▶ For some tasks, the word types are not that important, what matters is their **meaning**:
  - ▶ web search
  - ▶ Google *Ad Sense* (automatic ad placement)
- ▶ But what is a word meaning?

## Application 4: Word Meanings

- ▶ For some tasks, the word types are not that important, what matters is their **meaning**:
  - ▶ web search
  - ▶ Google *Ad Sense* (automatic ad placement)
- ▶ But what is a word meaning?

### Example: Dictionary Approach

#### **dog**

- ▶ is a **mammal**,
- ▶ descended from **wolf**,
- ▶ is commonly a **pet**,
- ▶ subtypes are **poodle**, **bulldog**, ...
- ▶ has **fur**,
- ▶ ...

# Why the Dictionary Approach is Problematic

- ▶ Such dictionaries have been tried for computers.  
e.g. WordNet
- ▶ They must be created by hand, which is a big problem:
  - ▶ expensive
  - ▶ only available for some languages
  - ▶ many new words missing
- ▶ We need dictionaries that can be generated automatically.

# Meaning as Word Use

- ▶ The philosopher **Ludwig Wittgenstein** said that a word's meaning is its use.

## Computational Counterpart

A word's meaning is given by how often it occurs together with other words.



# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-			
cat		-		
bark			-	
run				-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2		
cat		-		
bark			-	
run				-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	
cat		-		
bark			-	
run				-



# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat		-		
bark			-	
run				-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat	2	-		
bark			-	
run				-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	
bark			-	
run				-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	2
bark			-	
run				-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	2
bark	2		-	
run				-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	2
bark	2	1	-	
run				-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	2
bark	2	1	-	0
run				-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	2
bark	2	1	-	0
run	1			-



# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	2
bark	2	1	-	0
run	1	2		-

# Counting Tokens for Word Meaning

**Step 1:** Record in how many sentences words **occur together**

## Example

*The dog barked at the cat. The cat ran away. The dog ran after the cat. The dog kept barking. He also kept running.*

	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	2
bark	2	1	-	0
run	1	2	0	-

# From Vector to Vector Spaces

**Step 2:** Construct an  $n$ -dimensional vector space.

$n$  is given by the number of word types in the text

## 2-Dimensional Vector Space with *dog* and *cat*



	dog	cat	bark	run
dog	-	2	2	1
cat	2	-	1	2
bark	2	1	-	0
run	1	2	0	-

► *bark* more closely related to *dog*

► *run* more closely related to *cat*

# From Word Meaning to Text Meaning

**Step 3:** construct vector for whole text from word vectors

Meaning for *He either barks or runs and barks*



# Is that Realistic?

## ► Possible Concerns

- Is word meaning really just a bunch of numbers?
- In a real-word model, the vector space will have thousands of dimensions.
- But this might actually be what you have stored in your brain!

## Psycholinguistic Experiment

- Word association task
- The more similar the meaning vectors of two words, the faster test subjects recognize them as related.
- Also: Masked priming effects

# Usage for Web Search and Ad Sense

## ► Web Search

- construct meaning vector for every website
- convert search string to vector
- rank websites by vector similarity

## ► Ad Sense

- construct meaning vector for every website
- associate every ad with a vector
- pick ad that most closely matches website vector

## Bells and Whistles

- weigh words in search string by linear order  
aliens conspiracy  $\neq$  conspiracy aliens
- prioritize certain words in text
- compress vector space for efficiency

- ▶ For word meaning, the approach seems to work.
- ▶ But it is bad for sentence/text meaning.

## Example

The following two sentences receive the same vector:

- (1) a. Dog bites man!
- b. Man bites dog!

## A Friendly Pointer

How meaning actually works: **LIN 346 Language and Meaning**

# An Observation on Frequencies: Zipf's Law

- ▶ Word models care about word frequency.
- ▶ But there is a problem. . .

## Zipf's Law

The frequency of a type is inversely proportional to its rank.



## In Plain English

The most frequent word is

- ▶ 2 times as common as the 2nd most frequent word,
- ▶ 3 times as common as the 3rd most frequent word,
- ▶ and so on.



# An Observation on Frequencies: Zipf's Law

- ▶ Word models care about word frequency.
- ▶ But there is a problem. . .

## Zipf's Law

The frequency of a type is inversely proportional to its rank.




## In Plain English


The most frequent word is

- ▶ 2 times as common as the 2nd most frequent word,
- ▶ 3 times as common as the 3rd most frequent word,
- ▶ and so on.

# An Example from...the NBA?




# Visualizing Zipf Distributions



# Zipf's Law is Everywhere. . .

- ▶ A distribution is probably Zipfian if
  - ▶ there is a **long neck**:  
a few types make up the majority of tokens,
  - ▶ there is a **long tail**:  
most types only have 1 token (**hapax legomenon**)
- ▶ Surprisingly, Zipf's Law shows up in tons of places:
  - ▶ size of large cities in a country
  - ▶ citations for academic papers
  - ▶ frequencies of last names
  - ▶ frequencies of weekdays in text

## ... Even in Language!



# Stop Words

- ▶ About 150 words make up 50% of all English texts:  
*the, of, and, a, ...*
- ▶ These are called **stop words**.
- ▶ Stop words are not very informative for many applications.
- ▶ So they are usually discarded after the tokenization step.
- ▶ Failure to do so can greatly reduce the model's performance.

## Steps of Word Counting Model (Revised)

- 1 collect corpus
- 2 remove stop words
- 3 tokenize strings
- 4 count tokens for each type

# Stop Words

- ▶ About 150 words make up 50% of all English texts:  
*the, of, and, a, ...*
- ▶ These are called **stop words**.
- ▶ Stop words are not very informative for many applications.
- ▶ So they are usually discarded after the tokenization step.
- ▶ Failure to do so can greatly reduce the model's performance.

## Steps of Word Counting Model (Revised)

- 1 collect corpus
- 2 remove stop words
- 3 tokenize strings
- 4 count tokens for each type

## Example: A Text Without (Non)-Stop Words

- ▶ Stop words are much less informative than non-stop words.
- ▶ Just check the example below.

### **Stop Words only**

The                    having no                    on the



## Example: A Text Without (Non)-Stop Words

- ▶ Stop words are much less informative than non-stop words.
- ▶ Just check the example below.

### **Stop Words and Non-Stop Words**

The sun shone having no alternative on the nothing new

## Example: A Text Without (Non)-Stop Words

- ▶ Stop words are much less informative than non-stop words.
- ▶ Just check the example below.

### **Non-Stop Words only**

sun shone

alternative

nothing new

# An Important Consequence of Zipf's Law

- ▶ Texts mostly consist of stop words.
- ▶ Hence it can be difficult to get representative counts for non-stop words.

## Sparse Data Problem

- ▶ Most of the data is not informative.
- ▶ You need tons of data to have enough useful data.

# An Important Consequence of Zipf's Law

- ▶ Texts mostly consist of stop words.
- ▶ Hence it can be difficult to get representative counts for non-stop words.

## Sparse Data Problem

- ▶ Most of the data is not informative.
- ▶ You need tons of data to have enough useful data.

## Example

- ▶ Most models require corpora with at least a few million sentences.
- ▶ Really good models (e.g. Google translate) use billions of data points.

- ▶ Unigram (= word-based) models are surprisingly useful.
  - ▶ culturomics
  - ▶ stylistic analysis
  - ▶ authorship attribution
  - ▶ word meaning
  - ▶ web search
  - ▶ ad sense
  - ▶ and much more
- ▶ They just count words and do some number crunching with the frequencies.
- ▶ These models are everywhere, but they are also very simplistic.
- ▶ Quality depends largely on how much data is available.